

THE DESIGN PRESSURE

A dense Transformer evaluates every feed-forward unit for every token; activation sparsity is the counter-proposal. But Transformer sparsity is already emergent [3], needs a hard rectifier for exact zeros [4], and is input-dependent in *which* units fire, not merely how many [5]. Sparsity alone distinguishes nothing; what an architecture's sparsity depends on does.

WHAT BDH CHANGES

Dragon Hatchling (BDH) [1] makes it structural. In BDH-GPU, which Pathway open-sources, every layer multiplies two rectified (ReLU) vectors elementwise, and that product can never be denser than either. *Active* means a strictly non-zero entry there; a rectifier emits no negatives, so an active neuron is a positive amount of something rather than a signed coordinate. That product is the paper's $y_{t,l}$: Eq. (8) defines it as $(D_y \text{LN}(px))_+ \odot x$, Appendix E reads `relu(self.ln(a) @ decoder_y) * x`, and Fig. 14 counts its non-zeros. Naming collision: the code's `y_sparse` is the paper's rectified *intermediate*; the paper's *y* is the code's `xy_sparse`. Attention is the second change: query and key are one tensor, so the score matrix is a causally-masked Gram matrix of activations with **no softmax**. Positions whose neurons fired alike bind.

THE CLAIM, AND THE CORRECTION IT NEEDED

Section 6.4 reports "fewer neurons are active ... for more predictable input signals": at layer 2 of a 65,536-neuron model, 4.0–7.5% active while memorising a new word against about 2.5% while repeating it, a 1.6–3.0× ratio [1, Fig. 14]. We reproduced that protocol on three models trained here from the public code, all for 2,309 steps; it survives a 32-fold shrink: **1.47×, 2.12× and 1.87×** at 2,048, 8,192 and 16,384 neurons, each 1,280 sequences over five independently seeded samples, ±0.005 half-width. **In aggregate our three models agree with the paper: predictable text is quieter.** This model goes quiet on words it has just picked up from the text in front of it; it does not go quiet just because text is easy to guess. From the outside those two look identical, and they are not the same thing.

The same measurement narrows what that can mean. Two blocks of one sequence are predicted almost equally well yet differ 2.65-fold in activity, over 1,280 sequences of the 8,192-neuron model: the 13-letter warm-up runs at **9.65%** of layer 2 active at 0.0004 nats of mean surprise, the 56 repeated letters at **3.64%** at 0.004 nats. The warm-up is identical in every training sequence, so it lives in the **weights**; the repeated word is new each time, so it can only come from **context**. What *covaries* with the difference is not predictability but **where the knowledge came from** — a correlation, not an established cause. A reader tests this in a minute: inject an unpredictable letter mid-repetition and activity rises there, while the letter before is unchanged to nine decimals.

System	Where its sparsity comes from	Carried per token	Evidence status
Dense Transformer feed-forward	Emergent in training [3]; needs a hard rectifier for exact zeros [4]	A key-value cache, growing with context	Extensive, independent
Linear attention / SSM (Mamba [6])	None by design; dense recurrent state updates	A fixed-size recurrent state	Extensive, independent
BDH-GPU , measured here	Structural: elementwise product of two rectified vectors, never denser than either	A fixed-size synaptic state	Reproduced here at toy scale; author-reported at 65,536 neurons
BDH-CQ	Not public	Not public	Developer-reported only

EVIDENCE, LABELLED

REPRODUCED HERE: the three ratios, the counterexample, the surprise-injection check, browser-against-PyTorch parity (logit max difference 6.1e-5, twelve of twelve sparsity series bit-identical) — an implementation check, not architectural evidence. **DEVELOPER-REPORTED, NOT REPRODUCED:** the Fig. 14 band; 97.4% on Sudoku Extreme, which Pathway says does not reproduce from the open code; the 1B-to-600B scaling runs; BDH-CQ's ARC-AGI-1 result [2]. **NOT MODELLED:** BDH-CQ's internals, "proprietary" by its authors' own statement [2]; **it has no role here**, and inventing one would be worse. **WHERE IT IS LIKELY WORSE:** sparse non-negative activations plausibly cost capacity against dense superposition; no published comparison runs past GPT-2 scale.

LIMITATIONS AND MATURITY

Depth matters: layer 0 runs *backwards* at all three sizes, ratios below 1 (0.80, 0.80, 0.86): *more* neurons fire while repeating; layer 1 is positive at every size (1.22, 1.41, 1.33) and layer 3 is flat below 16,384. Scaling is **not monotonic** and we published the opposite before the third model finished. The wall-clock confound is gone, the short two retrained to 2,309 steps, the dip surviving at **40×** its half-width; none of the three published models finished the 4,000-step schedule, though separate runs that did finish move layer 2 only 1.4705 to 1.4620 at n=2,048, and 2.1201 to 2.2000 at n=8,192, a rise of 0.0799 against that run's own ±0.0086 spread. These are toys, 0.4M to 6.3M parameters on a synthetic task, as the paper's is. Surprise and sparsity do *not* track per letter (Pearson 0.30 across phases, Spearman -0.05 letter by letter: the relation is between blocks, not letters), killing a stronger claim. Most importantly: signature, not cause. We intervened twice. Fixing the word collapses n=2,048's 2.35× warm-up-to-repeat gap to 1.04×, but that control is degenerate and separates nothing. A 256-word pool does separate them, letters still in the weights, identity from context: that model is **4.5× easier by loss (0.038 against 0.174) yet shows a larger effect, 2.54× against 2.35×**, so the effect is not a monotone function of difficulty; K=16 (0.92) falls on the other side of the same line, unexplained. **MATURITY:** credible and outsider-checkable for this narrow property; the wider programme is earlier, its headline results developer-reported.

[1] Kosowski, Uznański, Chorowski, Stamirowska, Bartoszkiewicz, *The Dragon Hatchling*, arXiv:2509.26507 (2025), §6.4 and Fig. 14. [2] Engdahl, Kosowski, Chorowski et al., *BDH-CQ*, arXiv:2608.09888 (2026). [3] Li, You, Bhojanapalli et al., *The Lazy Neuron Phenomenon*, arXiv:2210.06313 (2022), ICLR 2023. [4] Mirzadeh, Alizadeh, Mehta et al., *ReLU Strikes Back*, arXiv:2310.04564 (2023). [5] Liu, Wang, Dao et al., *Deja Vu*, arXiv:2310.17157 (2023), ICML 2023. [6] Gu, Dao, *Mamba*, arXiv:2312.00752 (2023). **Where to continue, in order.** [1] §6.4 and Fig. 14 for the result reproduced here, then [1] §3.2 and Eq. (8) for the equations defining what is counted, then [3] for why sparsity alone distinguishes no architecture, then the artifact for the counterexample, which no paper states. — Reproduce the n=2,048 row: `python measure.py --ckpt checkpoints/bdh_n2048.pt --embd 64 --mult 32 --repeats 5`. Code and data MIT; bdh.py is Pathway's, unmodified, MIT.